
NicePiano

Cloud Computing 2020/2021

Matteo Orsini 1795119
Fabrizio Rossi 1815023




Main Features



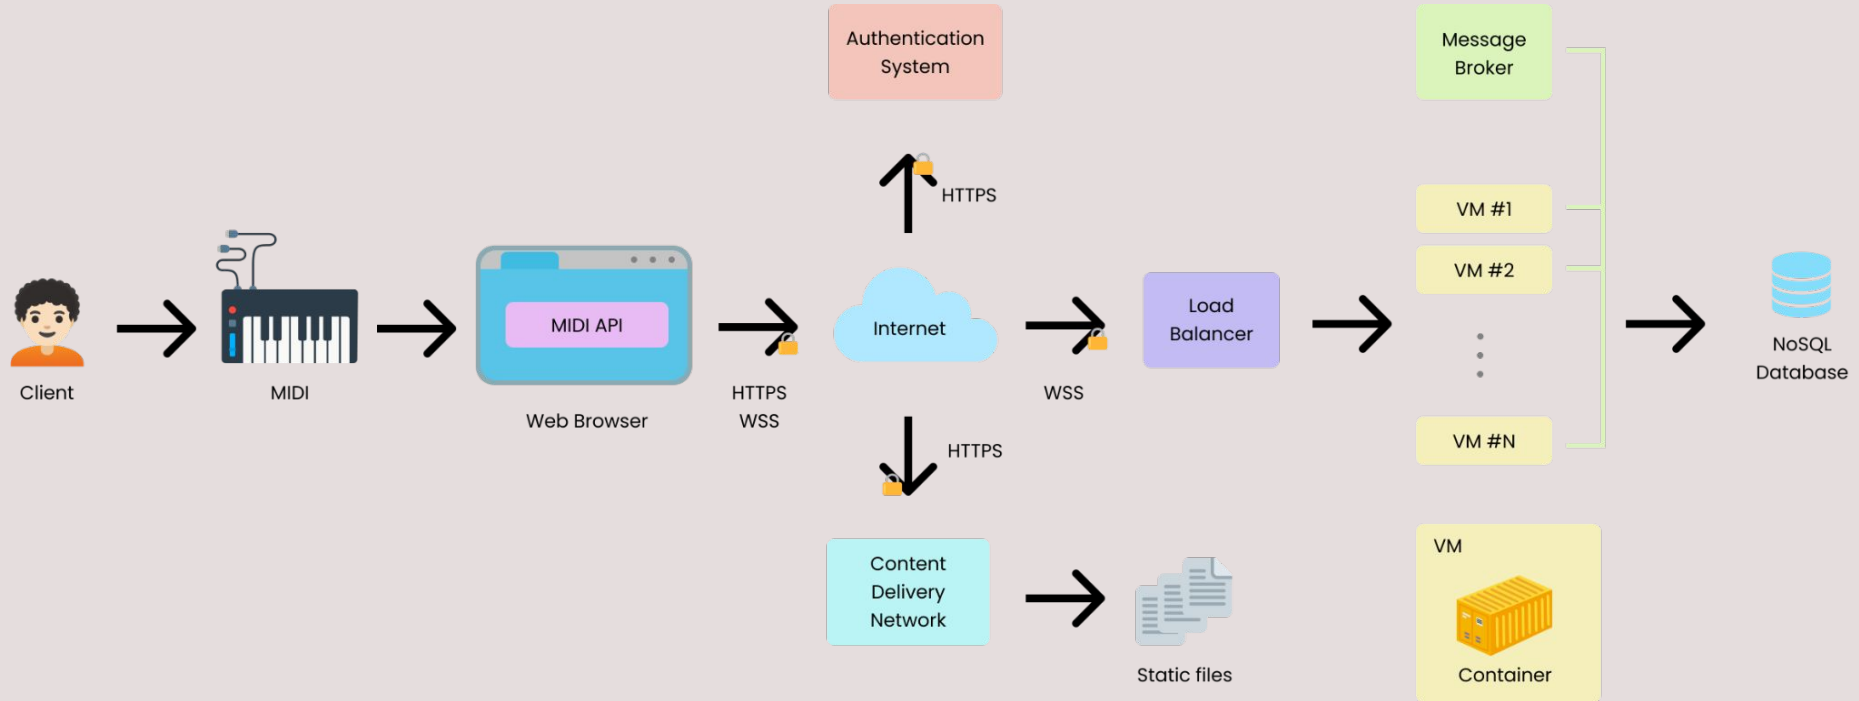
- **Web Application** based on cloud services
- Users can **play piano together**
 - In real time, using Web Sockets
- Virtual sessions called “**rooms**”
 - Notes played by one user will only be heard by users in the same room

User Interface

- Pressed keys visualization
- **Note visualization**
 - Scrolling rectangles
- Performance recording
- **Playback and looping**



Design - Architecture



Design - Server logic

- Receive events from Web Socket
- Broadcast them to users in the **same room**
 - **Problem:** possible only if users connected to same instance
 - **Solution:** message broker to share events
- Collect **latency** measurements
- Handle **recordings** operations with DB



Implementation - Front end

- Javascript for front end and back end
- React framework
- MIDI API (Chrome, Edge, Opera)
 - Tone.js and @tonejs/Piano
- aws-amplify for Amazon Cognito




Implementation - Back end

- Node.js runtime
- Docker engine
 - Dockerfile (e.g. DynamoDB table, port)
- Web Sockets (full-duplex, low latency)
 - **Socket.IO** (rooms)
- Redis for message exchanges
 - In-memory key-value storage
- Latency measurements (ping-pong)
- **aws-sdk** for DynamoDB




Deployment

- **Amazon Web Services (AWS)** with Educate account
- **Amazon Cognito** for authentication
 - User pool, client app ID
- **DynamoDB**
 - Recording table, on-demand capacity



Deployment - Front end

- Planned CI/CD with **CodePipeline** and **CodeBuild** (and S3)
 - Not available with Educate
- Front end needs to be served in **HTTPS** for MIDI
- **CloudFront** CDN
 - Generate certificate from ACM
 - Again both not available 😞
- Fallback: self hosted with self signed certificate




Deployment - Back end

- CI/CD with **CodePipeline**
- **ElasticBeanstalk**
 - Deploy web applications also using Docker engine
 - 1-3 EC2 t2.micro instances
 - Automatic scaling policy
 - Add instance if *Avg CPUUsage* > 66% for one minute
 - Remove instance if *Avg CPUUsage* < 20% for one minute
 - Load balancer
 - HTTPS: Custom signed certificate uploaded via IAM




Deployment - Back end cont.

- Message broker
 - **AWS ElastiCache** - fully managed Redis
 - Listed but not available
 - Confirmed by AWS support
 - Self hosted for tests




Validation

- Custom Node.js script
- Emulates **user behavior**
- Each user sends **random events** every 500ms:
 - 90% Note events
 - 10% Database requests



Validation cont.

- Ramp up: 5 users every 10s
 - Max 175 users
- Maintain load for 10 mins
- Ramp down





Validation cont.



- **More complex** since we used Web Sockets
- **Problem:** when second instance starts, load not spread
 - First instance clog up due to **warm-up** and web socket **stickiness** (only first connection distributed)
- **Solution:** disconnect probability
 - Every second 1% chance to disconnect and reconnect
 - New connection handled by load balancer
- **CloudWatch** to collect metrics (.csv format)
- Enable EC2 **detailed monitoring** for more precision (5m \Rightarrow 1m)



Results

- **6:00:** 100% CPU usage
 - New instance launched
- **9:00:** New instance **ready**
 - Start **load balancing** (disconnect probability)
- **16:00:** Ramp down
- **19:00:** Scale down
- **21:00:** Instance **shutdown**



Results cont.

- Response time = seconds between request to LB and response from VM
- **Good** response time even near 100% CPU usage
- Always < **600ms**
- DynamoDB **automatic scaling** with load in order to save money





Thanks for the attention!